

# HTTP: the language of web communication

Claudia Hauff

[cse1500-ewi@tudelft.nl](mailto:cse1500-ewi@tudelft.nl)



# Web technology overview

1. **HTTP: the language of web communication**
2. **HTML** & web app design
3. **JavaScript**: interactions in the browser
4. **Node.js**: JavaScript on the server
5. **CSS**: adding style
6. Node.js: advanced topics
7. Cookies & sessions
8. Web security

<https://vimeo.com/110256895>

# Learning goals

- **Describe** how Web servers and clients interact with each other.
- **Request** resources from web servers and understand the responses.
- **Describe** the different URL components.
- **Understand** and **employ** basic HTTP authentication.
- **Explain** the difference between HTTP and HTTPS.

# World Wide Web

vs.

# Internet

# The Web: a brief history

**World Wide Web:** a global system of interconnected hypertext documents available via the Internet  
(envisioned already in 1945)



- **1960s:** Precursor to the Internet (ARPANET) devised by the US department of Defense
  - Initial services: electronic mail, file transfer
- Late **1980s:** Internet opened to commercial interests
- **1989:** WWW created by Tim Berners-Lee (CERN)



#<<<>>> Lynx (web browser) - Wikipedia, the free encyclopedia (p1

#copyright

Your continued donations keep Wikipedia running!

## Lynx (web browser)

From Wikipedia, the free encyclopedia

Jump to: [navigation](#), [search](#)

CAPTION: Lynx

### Wikipedia Main Page displayed in Lynx

Wikipedia Main Page displayed in Lynx

Maintainer: Thomas Dickey  
Stable release: 2.8.5 (February 4, 2004) [\[\[+/-\]\]](#)  
Preview release: 2.8.6 (?) [\[\[+/-\]\]](#)  
OS: Cross-platform  
Use: web browser  
License: GPL  
Website: [lynx.isc.org](http://lynx.isc.org)

**Lynx** is a text-only [Web browser](#) and [Internet Gopher](#) client for use on cursor-addressable, character-cell [terminals](#).

Browsing in Lynx consists of highlighting the chosen link using cursor keys, or having all links on a page numbered and entering the chosen link's number. Current versions support [SSL](#) and many [HTML](#) features. Tables are linearized (scrunched together one cell after another without tabular structure), while frames are identified by name and can be explored as if they were separate pages.

Lynx is a product of the Distributed Computing Group within Academic Computing Services of the [University of Kansas](#), and was initially developed in 1992 by a team of students at the university (Lou Montulli, Michael Grobe and Charles Rezac) as a hypertext browser used solely to distribute campus information as part of a [Campus-Wide Information Server](#). In 1993 Montulli added an Internet interface and released a new version (2.0) of the browser [\[1\]](#) [\[2\]](#) [\[3\]](#).

# The Web: a brief history

**World Wide Web:** a global system of interconnected hypertext documents available via the Internet  
(envisioned already in 1945)



- **1960s:** Precursor to the Internet (ARPANET) devised by the US department of Defense
  - Initial services: electronic mail, file transfer
- Late **1980s:** Internet opened to commercial interests
- **1989:** WWW created by Tim Berners-Lee (CERN)
- **1994:** Netscape released its first Web browser







Location: about:

[Guided Tour](#) [What's New](#) [Questions](#) [Net Search](#) [Net Directory](#) [Newsgroups](#)

## Mosaic Netscape version 0.9 beta

Copyright © 1994 Mosaic Communications Corporation,  
All rights reserved.

This is *BETA* software subject to the license agreement set forth in the README file.  
Please read and agree to all terms before using this software.

Report any problems to [win\\_cbug@mcom.com](mailto:win_cbug@mcom.com).



Mosaic Communications, Mosaic Netscape, and the Mosaic Communications logo are trademarks of Mosaic Communications Corporation.

Any provision of Mosaic Software to the U.S. Government is with "Restricted rights" as follows: Use, duplication or disclosure by the Government is subject to restrictions set forth in subparagraphs (a) through (d) of the Commercial Computer Restricted Rights clause at FAR 52.227-19 when applicable, or in subparagraph (c) (1) (ii) of the Rights in Technical Data and Computer Software clause at DFARS 252.227-7013, and in similar clauses in the NASA FAR Supplement. Contractor/manufacturer is Mosaic Communications Corporation, 650 Castro Street, Suite 500, Mountain View, California, 94041.

# The Web: a brief history

**World Wide Web:** a global system of interconnected hypertext documents available via the Internet  
(envisioned already in 1945)



- **1960s:** Precursor to the Internet (ARPANET) devised by the US department of Defense
  - Initial services: electronic mail, file transfer
- Late **1980s:** Internet opened to commercial interests
- **1989:** WWW created by Tim Berners-Lee (CERN)
- **1994:** Netscape released its first Web browser
- **1995:** Microsoft released Internet Explorer v1
- **1998:** Google was founded
- **2002:** Mozilla released Firefox v1

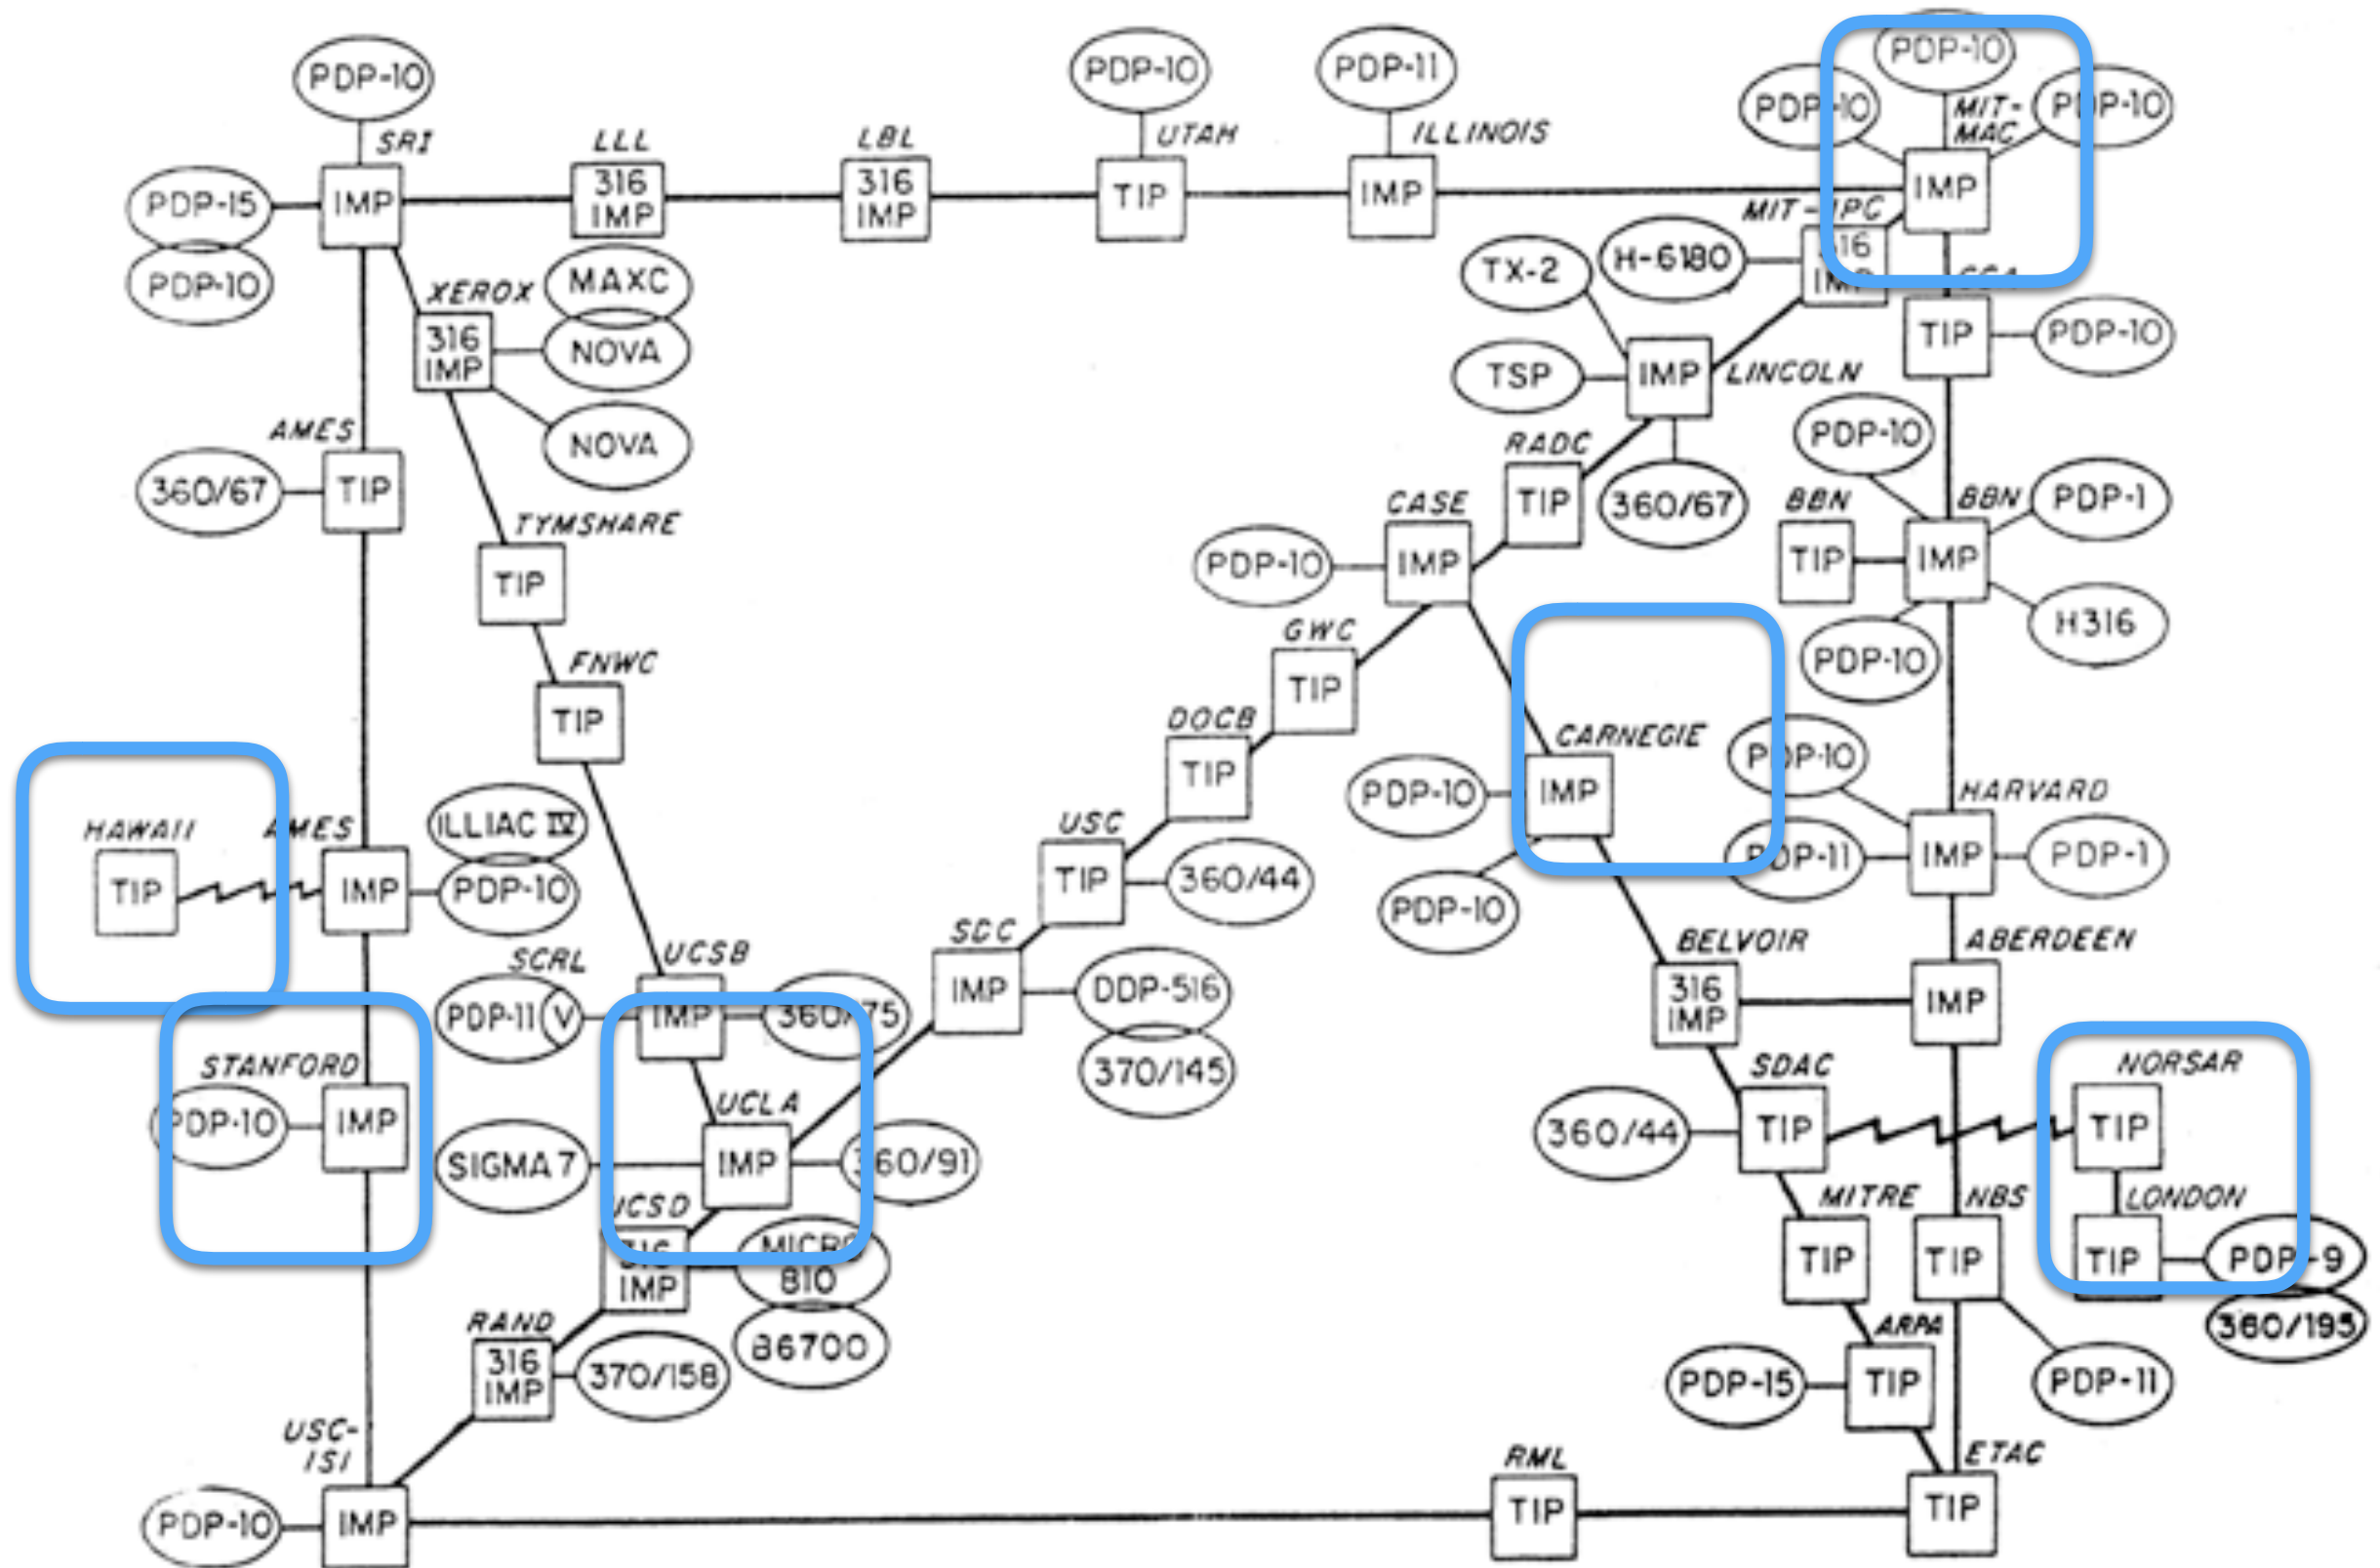


# Key aspects of the Internet

**Internet:** interconnected computer networks that span the globe; communicating through a common standard (TCP/IP)

- Sub-networks function **autonomously**
- **No centralised** control
- Devices **dynamically** join/leave the network
- Devices interact through **open standards**
- **Easy** to use: server/client software widely available

## ARPA NETWORK, LOGICAL MAP, SEPTEMBER 1973



## State of the Internet in 1973

Image source: <http://qz.com/860873/a-1973-map-of-the-internet-charted-by-darpa/>



# Two important organisations

- Internet Engineering Task Force (IETF)

“The mission of the IETF is to make the Internet work better by producing high quality, relevant technical documents that influence the way people design, use, and manage the Internet.”

**Request for Comments (RFC)**

- World Wide Web Consortium (W3C)

“The W3C mission is to lead the World Wide Web to its full potential by developing protocols and guidelines that ensure the long-term growth of the Web.”

# HTTP messages



# HTTP/1.1

**RFC 2068**

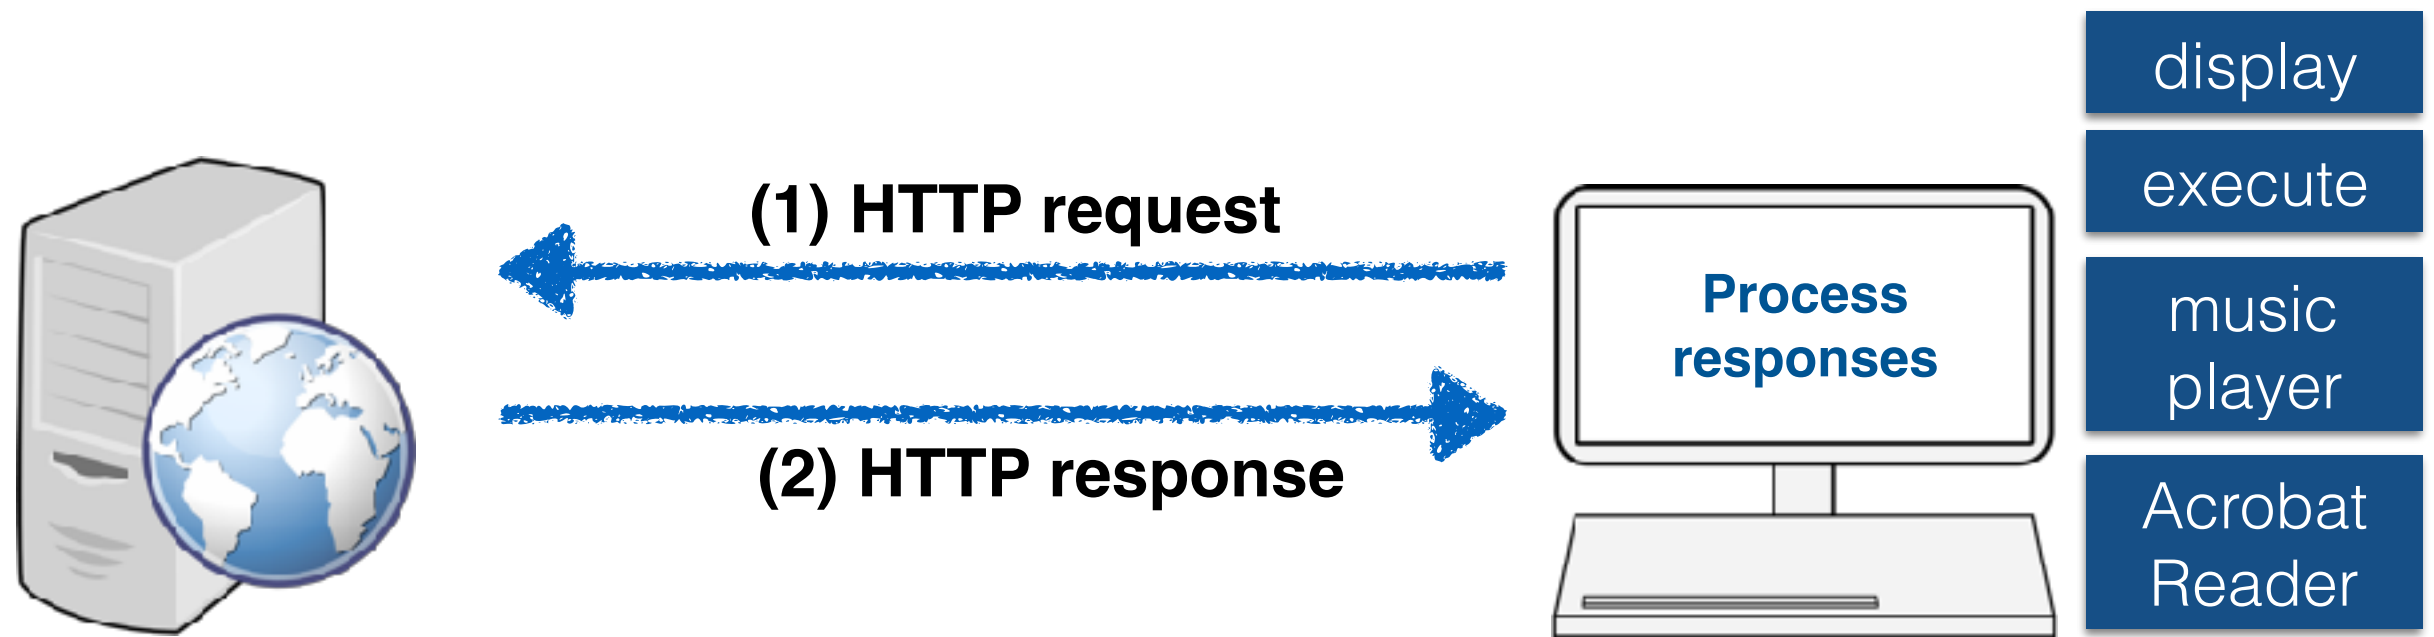
**1997**

# HTTP/2

**RFC 7540**

**2015**

# Web servers and clients



- Servers wait for data requests
- Answer thousands of clients simultaneously
- Host **web resources** (content with an identity)
- Clients are often browsers
- Telnet



# Network communication

- Conceptual model **Open Systems Interconnection** (OSI) from 1995
- Network protocols matched to layers
- Many network protocols exist, we care about three in particular:

<b>IP</b>	Internet Protocol
<b>TCP</b>	Transmission Control Protocol
<b>HTTP</b>	Hypertext Transfer Protocol

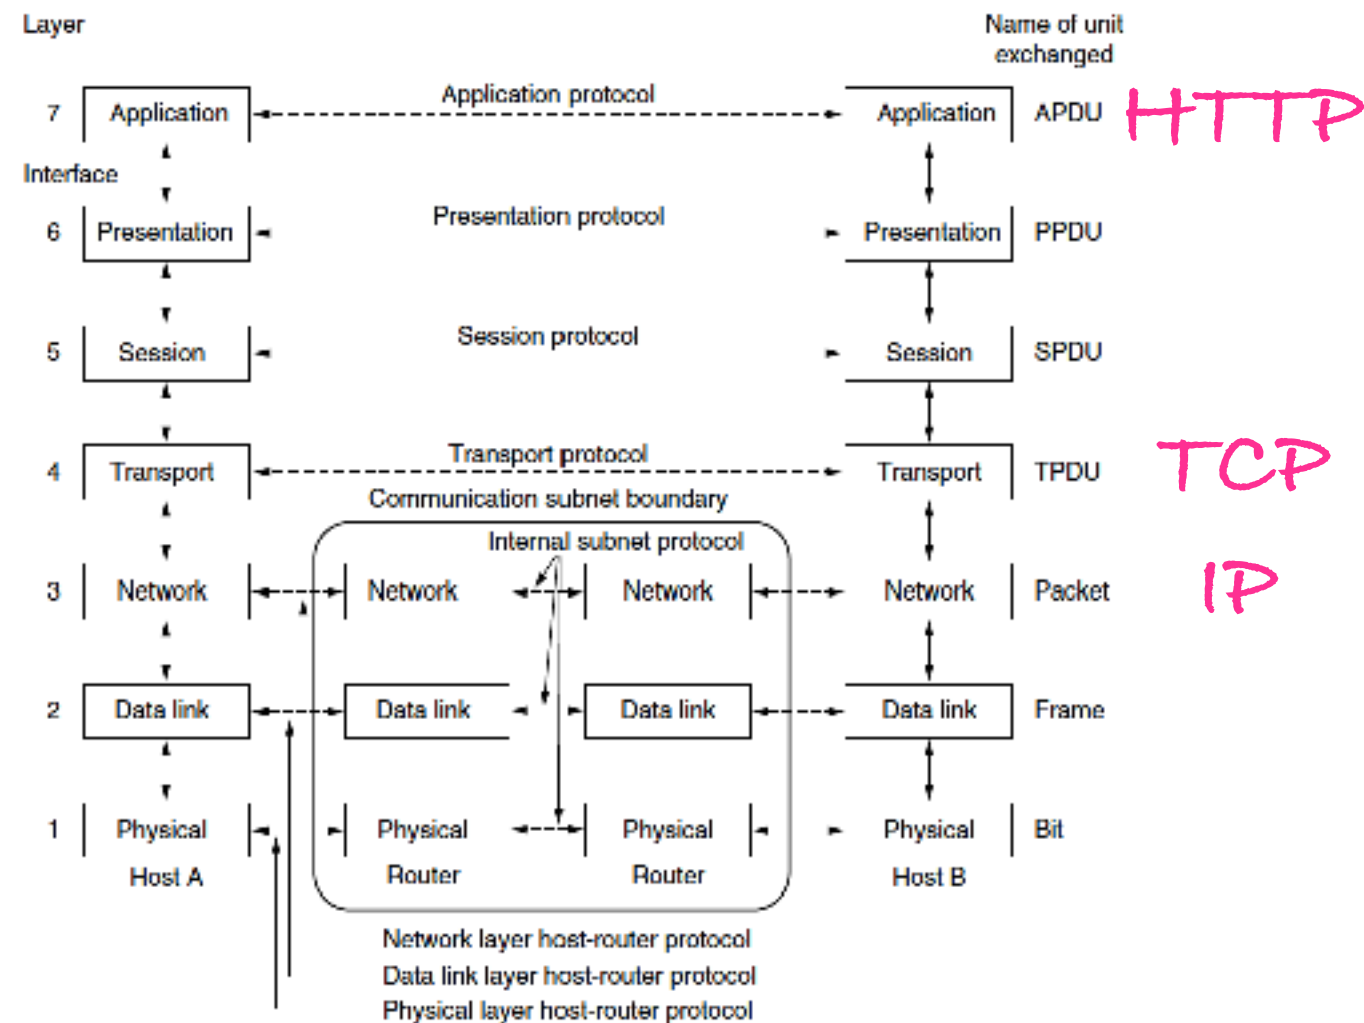


Image src: Computer Networks (5th edition), Tanenbaum & Whetherall, p. 42

HTTP uses **reliable** data-transmission protocols (inherited from TCP).

# HTTP request message

plain text, line-oriented character sequences

```
GET / HTTP/1.1
Host: www.tudelft.nl
User-Agent: Mozilla/5.0 (Macintosh; Intel Mac OS X
10.9; rv:31.0) Gecko/20100101 Firefox/31.0
Accept: text/html,application/xhtml+xml,application/
xml;q=0.9,*/*;q=0.8
Accept-Language: en-gb,en;q=0.5
Accept-Encoding: gzip, deflate
DNT: 1
Cookie:
__utma=1.20923577936111.16111.19805.2;utmcmd=(none);
```

# HTTP response message

```
HTTP/1.1 200 OK
```

start line

```
Date: Fri, 01 Aug 2014 13:35:55 GMT
```

```
Content-Type: text/html; charset=utf-8
```

```
Content-Length: 5994
```

```
Connection: keep-alive
```

header fields

```
Set-Cookie: fe_typo_user=d5e20a55a4a92e0;  
path=/; domain=tudelft.nl
```

name:value

```
[...]
```

```
Server: TU Delft Web Server
```

```
...
```

```
...
```

body  
(optional)

# HTTP headers dissected



# Well-known header fields

<b>Content-Type</b>	Entity type
<b>Content-Length</b>	Length/size of the message
Content-Language	Language of the entity sent (e.g. English)
<b>Content-Encoding</b>	Data transformations applied to the entity
Content-Location	Alternative location of the entity
Content-Range	Range defines the pieces sent for partial entities
<b>Content-MD5</b>	Checksum of the content
<b>Last-Modified</b>	Date on which this entity was created/modified
<b>Expires</b>	Date at which the entity will become stale
Allow	Lists the legal request methods for the entity
<b>Connection &amp; Upgrade</b>	Protocol upgrade

Entity bodies contain **raw** data. **Header** needed to **interpret** the data.

# Content-Type

- **MIME** types are attached to all HTTP object data

Multipurpose Internet **Mail** Extensions

historic reasons

- Type determines clients' reaction to the data
- Pattern: [primary object type]/[subtype]  
e.g. text/plain, text/html, image/jpeg,  
video/quicktime

# Content types are diverse

Most popular
text/html
image/jpeg
text/xml
application/rss+xml
text/plain
application/xml
text/calendar
application/pdf
application/atom+xml
unknown/unknown

Least popular
application/pgp-keys
application/x-httpd-php4
chemical/x-pdb
model/mesh
application/x-perl
audio/x_mpegurl
application/bib
application/postscript
application/x-msdos-program

# Content-Length

- Indicates the **size** of the entity body
- Necessary to detect premature message truncation (e.g. due to a server crash, faulty proxy)
- Essential for **persistent connections** to discover where one HTTP message ends and the next begins

**Persistent connections** reuse the same TCP connection for multiple HTTP request/response messages.



# Content-Encoding

- Commonly either `gzip`, `compress` (Unix compression), `deflate` (zlib compression) or `identity` (no encoding)
- Servers aim for **encodings** that clients understand
  - Clients send a list of acceptable encodings in a corresponding request header:  
`Accept-Encoding: gzip, deflate`
- **Compression** saves network bandwidth but increases processing costs


# Content-MD5

## Message Digest

- HTTP messages are sent via TCP/IP
- **However**: the Internet is huge, many servers interact to transport a message with different implementations (**bugs!**) of established protocols
- **Sanity check**: Sender generates a **MD5 checksum** of the content (hashed into a 128 bit value) to detect unintended modifications of the content
- Has been removed from the HTTP/1.1 specification (2014), but remains heavily in use

# Expires

- **Web caches** keep copies of *popular* resources



- Advantages of Web caches
  - A. **Reduction** of redundant data transfer
  - B. **Reduction** of network bottlenecks
  - C. **Reduction** demand on origin servers
  - D. **Reduced** distance delay
- **Expires** indicates when the resource is no longer valid

# Expires & Cache-Control

- Content on the origin server can change
- Caches need to ensure that their copies are **in sync** with the origin server
- Caches can revalidate their copies at any time (inefficient)
- **Expires** in HTTP response header indicates a **resource's expiration date** in **absolute** terms — date determines when the cache revalidates
- **Cache-Control** indicates a resource's expiration date in **relative** terms (seconds since being sent)

# Expires & Cache-Control

- Content on the origin server
- Caches need to ensure the origin server
- Caches can revalidate
- Expires in HTTP response **resource's expiration** determines when the cache
- Cache-Control: indicates **relative** terms (number

Headers Cookies Params Response Timings Stack Trace

Request URL: <https://www.theguardian.com/uk>  
Request method: GET  
Remote address: 151.101.37.111:443  
Status code: 200 ⓘ Edit and Resend Raw headers  
Version: HTTP/2.0

Filter headers

Response headers (1.511 KB)

- accept-ranges: bytes
- age: 74
- cache-control: max-age=60, stale-while-revali...tale-if-error=864000, private
- content-encoding: gzip
- content-length: 138026
- content-security-policy: default-src https;;data;; connect-src https: wss:
- content-type: text/html; charset=utf-8
- date: Fri, 19 Oct 2018 08:35:14 GMT
- etag: W/"hash-3272292653060572815"
- link: rel=preload; as=script; nopush
- set-cookie: GU\_mvt\_id=894690; expires=Thu,...th=/; domain=.theguardian.com
- set-cookie: GU\_geo\_continent=EU; path=/;
- strict-transport-security: max-age=31536000; includeSubDomains; preload
- vary: Accept-Encoding,User-Agent
- x-content-type-options: nosniff
- X-Firefox-Spdy: h2
- x-frame-options: SAMEORIGIN

# Last-Modified

- Contains the date when the resource was last **altered**
- **No indication** about the amount of changes
- Often used in combination with **If-Modified-Since** for cache revalidation requests: origin server only returns the document if it changed since the given date
- Last-Modified dates are **not reliable**



# Connection & Upgrade

- HTTP/1.1: the client **always** initiates the conversation

## Polling the 2018 Midterm Elections in Real Time

The Upshot has partnered with [Siena College](#) to poll dozens of the most competitive House and Senate races across the country. Our poll results are updated in real time, after every phone call. We hope to help you understand how polling works, and why it sometimes doesn't.

[Related article »](#)

[Sign up for alerts](#)

185,872  
calls made so far

We're c

Virginia

Can a Tea P

Continues so

The screenshot shows a Twitter interface. At the top, there are navigation links: Home (with a house icon), Notifications (with a bell icon), and Messages (with an envelope icon). A search bar is on the right. Below the navigation bar, there's a section for a user profile. The profile picture is a circular image of a woman with glasses. The name is Claudia Hauff, and the handle is @CharlotteHase. Below the name, there are three statistics: Tweets (3,800), Following (742), and Followers (1,611). To the right of the profile, there's a text input field with the placeholder "What's happening?". Below the input field, there's a button that says "See 1 new Tweet". Below the button, there's a tweet from TEDxAmsterdam (@TEDxAmsterdam) posted 3 minutes ago. The tweet text is: "And what a brilliant #TEDxAmsterdam Award pitch it was! Applications for the 2018 Award are open until 5 Oct. Make like @GoodShipping1 did and send us your idea. Who knows how far it will take you? #thebigx #thebigunknown".

Home Notifications Messages Search

What's happening?

See 1 new Tweet

**CLAUDIA HAUFF**  
@CharlotteHase

Tweets 3,800 Following 742 Followers 1,611

**TEDxAmsterdam** @TEDxAmsterdam · 3m

And what a brilliant #TEDxAmsterdam Award pitch it was! Applications for the 2018 Award are open until 5 Oct. Make like @GoodShipping1 did and send us your idea. Who knows how far it will take you? #thebigx #thebigunknown

# Connection & Upgrade

- HTTP/1.1: the client **always** initiates the conversation
- Simulating a **server-side push** of data:
  - **Polling**: client regularly sends HTTP requests
  - **Long polling**: client sends an HTTP request and the server holds it open until new data arrives
  - Issues: wasted bandwidth, complex backend
- From simulation to solution (standardized in 2011): the **WebSocket protocol** enables bidirectional communication between client & server

# Connection & Upgrade

## Lecture 4

The screenshot shows a web browser at `localhost:3000/play`. The main content area displays a grid of letters A through T, and a row of red hash symbols (`#####`) above four black circles. The browser's developer tools are open, showing the Network tab. The network log lists several requests, with the last one being a GET request to `/` via a websocket connection, which has a status code of 101. The headers for this request are visible, showing `Sec-WebSocket-Extensions: permessage-deflate`, `Sec-WebSocket-Key: PTXgPhQoo1j61/Zklj/uoQ==`, and `Sec-WebSocket-Version: 13`. The upgrade is confirmed to be `websocket`.

Status	Method	F...	D...	Cause	Type	Transferred	Size	0 ms	640 ms	1.2
200	GET	config.js	loc...script	js	970 B	864 B	→ 28 ms			
200	GET	status...	loc...script	js	1.29 KB	0.89 KB	→ 32 ms			
200	GET	messag...	loc...script	js	2.43 KB	2.13 KB	→ 35 ms			
200	GET	alphab...	loc...script	js	2.86 KB	2.56 KB	→ 41 ms			
200	GET	ui-right...	loc...script	js	1.59 KB	1.29 KB	→ 37 ms			
200	GET	interact...	loc...script	js	9.04 KB	8.74 KB	→ 37 ms			
200	GET	style.css	loc...stylesheet	css	608 B	316 B	→ 4 ms			
101	GET	/	loc...websocket	plain	129 B	0 B				

11 requests | 10.41 MB / 10.42 MB transferred | Finish: 2.10 s | DOMContentLoaded: 2.10 s | load: 2.14 s

Headers

Request URL: `http://localhost:3000/`  
Request method: GET  
Remote address: `[::1]:3000`  
Status code: **101** ⓘ Edit and Resend Raw headers  
Version: HTTP/1.1

Filter headers

- Pragma: no-cache
- Sec-WebSocket-Extensions: permessage-deflate
- Sec-WebSocket-Key: PTXgPhQoo1j61/Zklj/uoQ==
- Sec-WebSocket-Version: 13
- Upgrade: websocket

# Connection & Upgrade

- Client and server have to agree to the **protocol upgrade**
- **Client initiates** the upgrade with two request headers:

```
Connection:Upgrade  
Upgrade:[protocols]
```

- Server responds with a 101 **Switching Protocols** status if a protocol upgrade is possible
- Once established, both the client and server can push data

# Common status codes

## 10.4.3 402 Payment Required

This code is reserved for future use.

<b>1xx</b>	Informational	(101 Switching ...)
------------	---------------	---------------------

<b>2xx</b>	Success	(200 OK)
------------	---------	----------

<b>3xx</b>	Redirected	
------------	------------	--

<b>4xx</b>	Client error	(404 Not Found)
------------	--------------	-----------------

<b>5xx</b>	Server error	
------------	--------------	--

In practice only a few codes per category are supported

# HTTP methods

`GET / HTTP/1.1`

`Host: www.tudelft.nl`

`User-Agent: Mozilla/5.0 (Macintosh; Intel Mac OS X 10.9; rv:31.0) Gecko/20100101 Firefox/31.0`

`Accept: text/html,application/xhtml+xml;q=0.9,*/*;q=0.8`

`Accept-Language: en-gb,en;q=0.5`

`Accept-Encoding: gzip, deflate`



# Common HTTP methods

**GET**

Get a document from the Web server.

**HEAD**

Get the header of a document from the Web server.

**POST**

Send data from the client to the server for processing.

**PUT**

Save the body of the request on the server.

**TRACE**

Trace the message through proxy servers to the server.

**OPTIONS**

Determine what methods can operate on a server.

**DELETE**

Remove a document from a Web server.


Servers may implement more or fewer methods than shown.

```
2018-11-09 16:08:26 wlan-145-94-166-99 in ~  
[o → telnet microsoft.com 80  
Trying 40.112.72.205...  
Connected to microsoft.com.  
Escape character is '^]'.  
HEAD / HTTP/1.1  
host:microsoft.com
```

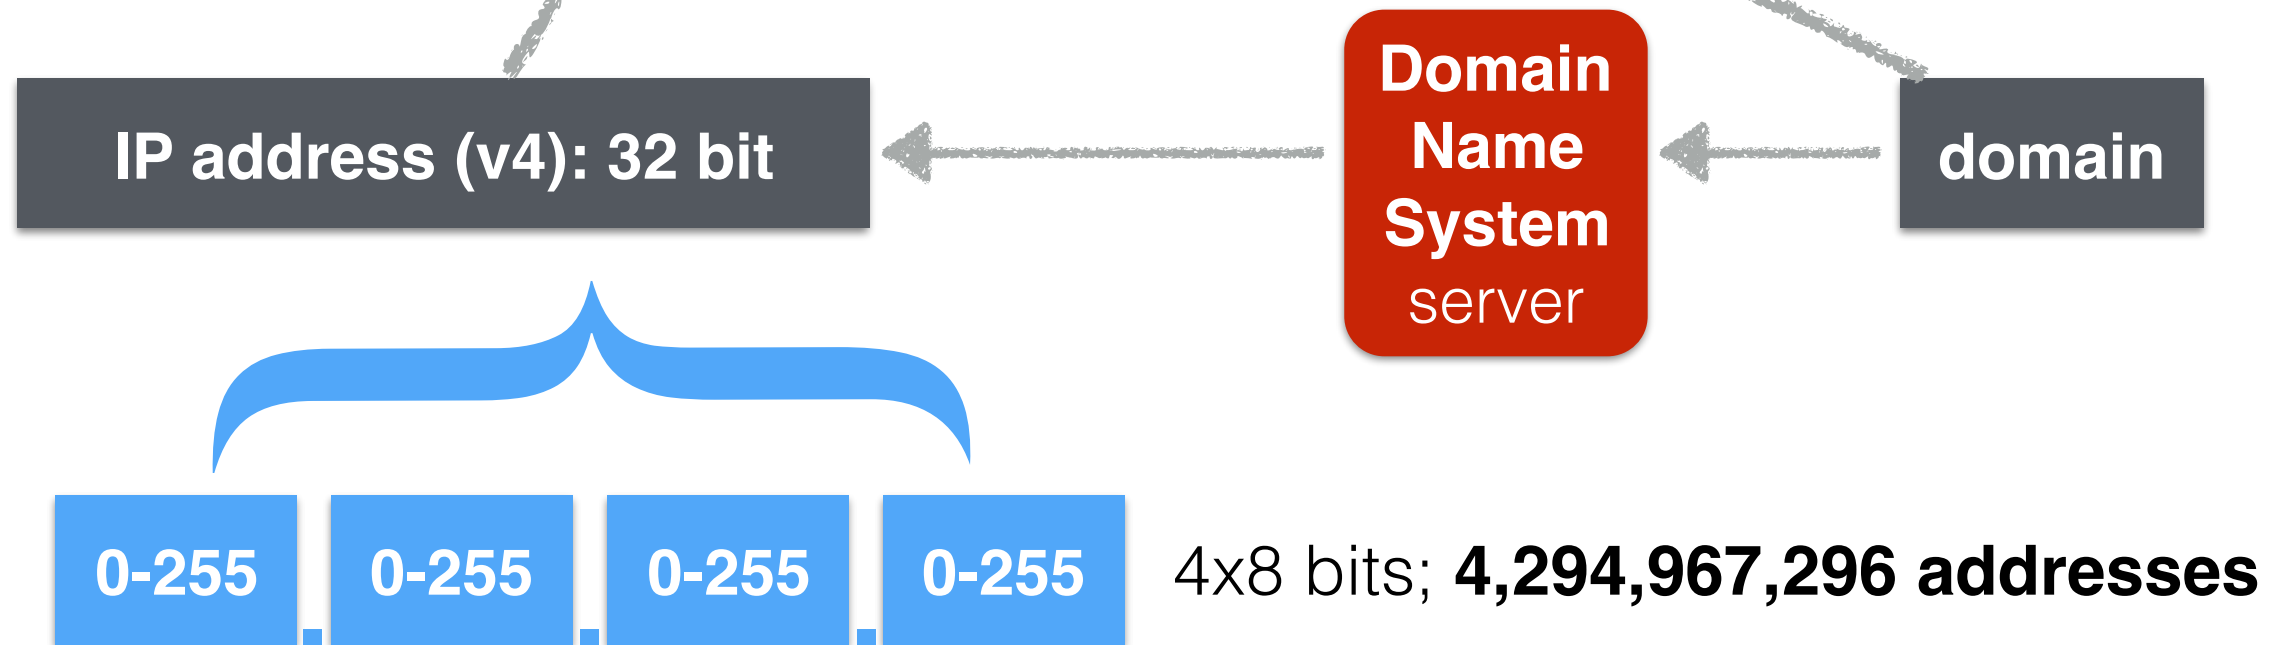
```
HTTP/1.1 301 Moved Permanently  
Date: Fri, 09 Nov 2018 15:08:45 GMT  
Server: Kestrel  
Location: https://www.microsoft.com/  
█
```

Telnet opens a **TCP connection** to a Web server; chars are typed directly into the port. The server treats telnet as web client, the returned data is displayed onscreen.

# From domain to IP address



```
2018-11-09 16:23:25 wlan-145-94-166-99 in ~  
[o → telnet microsoft.com 8  
Trying 40.112.72.205.  
Connected to microsoft.com.  
Escape character is '^]'.  
█
```




**IP address (v6): 128 bit**

2A41 ■ 0000 ■ 3341 ■ 3AAA ■ 0000 ■ FFFF ■ 86AA ■ B92E

# From domain to IP address

## IPv6 Adoption

We are continuously measuring the availability of IPv6 connectivity among Google users. The graph shows the percentage of users that access Google over IPv6.



# Uniform Resource Locators (URLs)

# URL syntax

- Uniform resource locators offer a standardised way to point to any resource on the Internet
- Not restricted to `http`, however, the syntax slightly varies from scheme to scheme
- General format:

```
<scheme>://<user>:<password>@<host>:<port>/<path>;<params>?<query>#<frag>
```



# URL syntax

- Uniform resource locators offer a standardised way to point to any resource on the Internet
- Not restricted to `http`, however, the syntax slightly varies from scheme to scheme
- General format:

**<scheme>**://<user>:<password>@<host>:<port>/<path>;<params>?<query>#<frag>

*determines the protocol to use when connecting to the server*

# URL syntax

- Uniform resource locators offer a standardised way to point to any resource on the Internet
- Not restricted to `http`, however, the syntax slightly varies from scheme to scheme
- General format:

`<scheme>://<user>:<password>@<host>:<port>/<path>;<params>?<query>#<frag>`

*the username/password (to access a protected resource)*

# URL syntax

- Uniform resource locators offer a standardised way to point to any resource on the Internet
- Not restricted to `http`, however, the syntax slightly varies from scheme to scheme
- General format:

`<scheme>://<user>:<password>@<host>:<port>/<path>;<params>?<query>#<frag>`

*the domain name or numeric IP address of the server*

# URL syntax

- Uniform resource locators offer a standardised way to point to any resource on the Internet
- Not restricted to `http`, however, the syntax slightly varies from scheme to scheme
- General format:

```
<scheme>://<user>:<password>@<host>:<port>/<path>;<params>?<query>#<frag>
```

*the port on which the server is expecting requests*

# URL syntax

- Uniform resource locators offer a standardised way to point to any resource on the Internet
- Not restricted to `http`, however, the syntax slightly varies from scheme to scheme
- General format:

`<scheme>://<user>:<password>@<host>:<port>/<path>;<params>?<query>#<frag>`

*the local path to the resource*

# URL syntax

- Uniform resource locators offer a standardised way to point to any resource on the Internet
- Not restricted to `http`, however, the syntax slightly varies from scheme to scheme
- General format:

```
<scheme>://<user>:<password>@<host>:<port>/<path>;<params>?<query>#<frag>
```

*additional input parameters applications may require*



# URL syntax

- Uniform resource locators offer a standardised way to point to any resource on the Internet
- Not restricted to `http`, however, the syntax slightly varies from scheme to scheme
- General format:

```
<scheme>://<user>:<password>@<host>:<port>/<path>;<params>?<query>#<frag>
```

*parameters passed to gateway resources (e.g. a search engine)*

# URL syntax

- Uniform resource locators offer a standardised way to point to any resource on the Internet
- Not restricted to `http`, however, the syntax slightly varies from scheme to scheme
- General format:

`<scheme>://<user>:<password>@<host>:<port>/<path>;<params>?<query>#<frag>`

*the name of a piece of a resource; only used by the client*

common convention: name1=value1&name2=value2&...


web images

🔍 Netherlands


**30 Hotels in Delft - Best Price Guarantee.** AD

Book your Hotel in Delft online. No reservation costs. Great rates.  
Book Now, No Booking Fees, Book for Tonight  
[www.booking.com/Delft/Hotels](https://www.booking.com/Delft/Hotels) Report Ad


**Your cityguide to Delft — Delft.com**

Discover the story of William of Orange, get up close and personal with Johannes Vermeer and see how the world-famous Delft Blue is made.  
 <https://www.delft.com>


**Delft 2018: Best of Delft, The Netherlands Tourism - TripAdvisor**



Vermeer's birthplace and a true gem, Delft sits between The Hague and Rotterdam in the country's southwest. The city's name comes from the Dutch word for digging, fitting since canals are a highlight here. Others include the 13th-century Old Church, the 15th-century New Church and the ...  
 [https://www.tripadvisor.com/Tourism-g188626-Delft\\_South\\_Holland\\_Provi...](https://www.tripadvisor.com/Tourism-g188626-Delft_South_Holland_Provi...)

**Delft - Wikipedia**

Delft (  [listen]) is a city and municipality in the province of South Holland, Netherlands. It is located between Rotterdam, to the southeast, and The Hague, to the northwest.  
 <https://en.wikipedia.org/wiki/Delft>

**Delft travel - Lonely Planet**

Explore Delft holidays and discover the best time and places to visit. | An amalgam of austere medieval magnificence and Golden Age glory, Delft's exquisite town centre is a hugely popular Dutch day-trip destination, awash with visitors strolling its narrow, canal-lined streets and central Markt.  
 <https://www.lonelyplanet.com/the-netherlands/the-randstad/delft>



**Delft**

delft.nl

Delft is a city and municipality in the province of South Holland, Netherlands. It is located between Rotterdam, to the southeast, and The Hague, to the northwest. Together with them, it is part of the Rotterdam–The Hague metropolitan area and the Randstad. [More at Wikipedia](#)

**Country:** Netherlands

**Body:** Municipal council

**Mayor:** Marja van Bijsterveldt (CDA)

**Spode Delft Plate - Free Shipping, Orders \$99 or More** AD

Free Shipping, Orders \$99 or More. Get Details and Your Pattern  
[Replacements.com/Spode](https://Replacements.com/Spode) Report Ad

**CV ketel Delft Voordeelactie - Welke korting ga jij kiezen?**

Alleen nog in Maart: €100 huurtegoed of gratis slimme thermos  
€199

# Schemes: more than just http

```
http://<host>:<port>/<path>?<query>#<frag>
```

```
https://<host>:<port>/<path>?<query>#<frag>
```

```
mailto:<valid-email-address>
```

```
file://<host>/<path>
```

```
file:///Users/my_home_dir/tmp.html
```

```
ftp://<user>:<passwd>@<host>:<port>/<path>;<params>
```

# Relative vs. absolute URLs

## base url

absolute

`http://www.st.ewi.tudelft.nl/~hauff/new/index.html`

relative

```
<h1>Visualizations</h1>
```

```
<ol>
```

```
  <li><a href="vis/trecvis.html">TREC</a></li>
```

```
  <li><a href=" ../airsvis.html">AIRS</a></li>
```

```
</ol>
```



URLs can be complex.  
RFC 3986 governs conversion rules.

`http://www.st.ewi.tudelft.nl/~hauff/new/vis/trecvis.html`

`http://www.st.ewi.tudelft.nl/~hauff/airsvis.html`

# URL design restrictions

- **Initial design goals:** *portable* across protocols and human *readable* (no invisible/non-printing chars.)
- URLs initially restricted to a very small “safe” alphabet: ASCII Heavily biased in favour of English speakers

latin alphabet, 0123456789 - \_ . ~  
and additional reserved chars like ! ( ) @ &

- **Added later:** character encoding, e.g. whitespace as %20

<http://правительство.рф>

# Punycode (RFC3492)

“Punycode is a simple and efficient transfer encoding syntax designed for use with Internationalized Domain Names in Applications. It **uniquely** and **reversibly** transforms a Unicode string into an ASCII string.”

`http://правительство.рф`

`=> http://:xn-80aealotwbjpid2k.xn-plai`

`http://nl.wikipedia.org/wiki/Itali%C3%AB`

`=> http://nl.wikipedia.org/wiki/Itali%C3%AB`

A potential security issue in *mixed scripts*: `http://paypal.com`



The background of the slide is split diagonally from the top-left corner to the bottom-right corner. The upper-left portion is a solid, vibrant pink, while the lower-right portion is a solid white.

# Authentication

# Authentication

So far: HTTP as **anonymous**, **stateless** request/response protocol. The same request, sent by different clients, is treated in exactly the same manner.

Now: identification via

- A. HTTP headers
- B. Client IP address tracking
- C. Fat URLs
- D. User login (HTTP Basic Authentication)

In lecture 7: Cookies & Sessions.

# User-related HTTP header fields

**From**

Request

User's email address

**mostly Web crawler**

**User-Agent**

Request

User's browser

**device customization**

**Referer**

Request

Page the user came from

**user interests**

**Client-IP**

Request  
(Extension)

Client's IP address

**Authorization**

Request

Username & password

# Client IP address tracking

- Idea: **Client IP address** as user identifier
- **Several issues:**
  - A. IP addresses describe the **machine**, not the user
  - B. ISPs **dynamically assigned IP** addresses to users
  - C. Users may access the Web through **firewalls**
  - D. HTTP **proxies** and **gateways** open new TCP connections (IP of the proxy/gateway is shown), **X-Forwarded-For** might help (it quickly gets complicated)

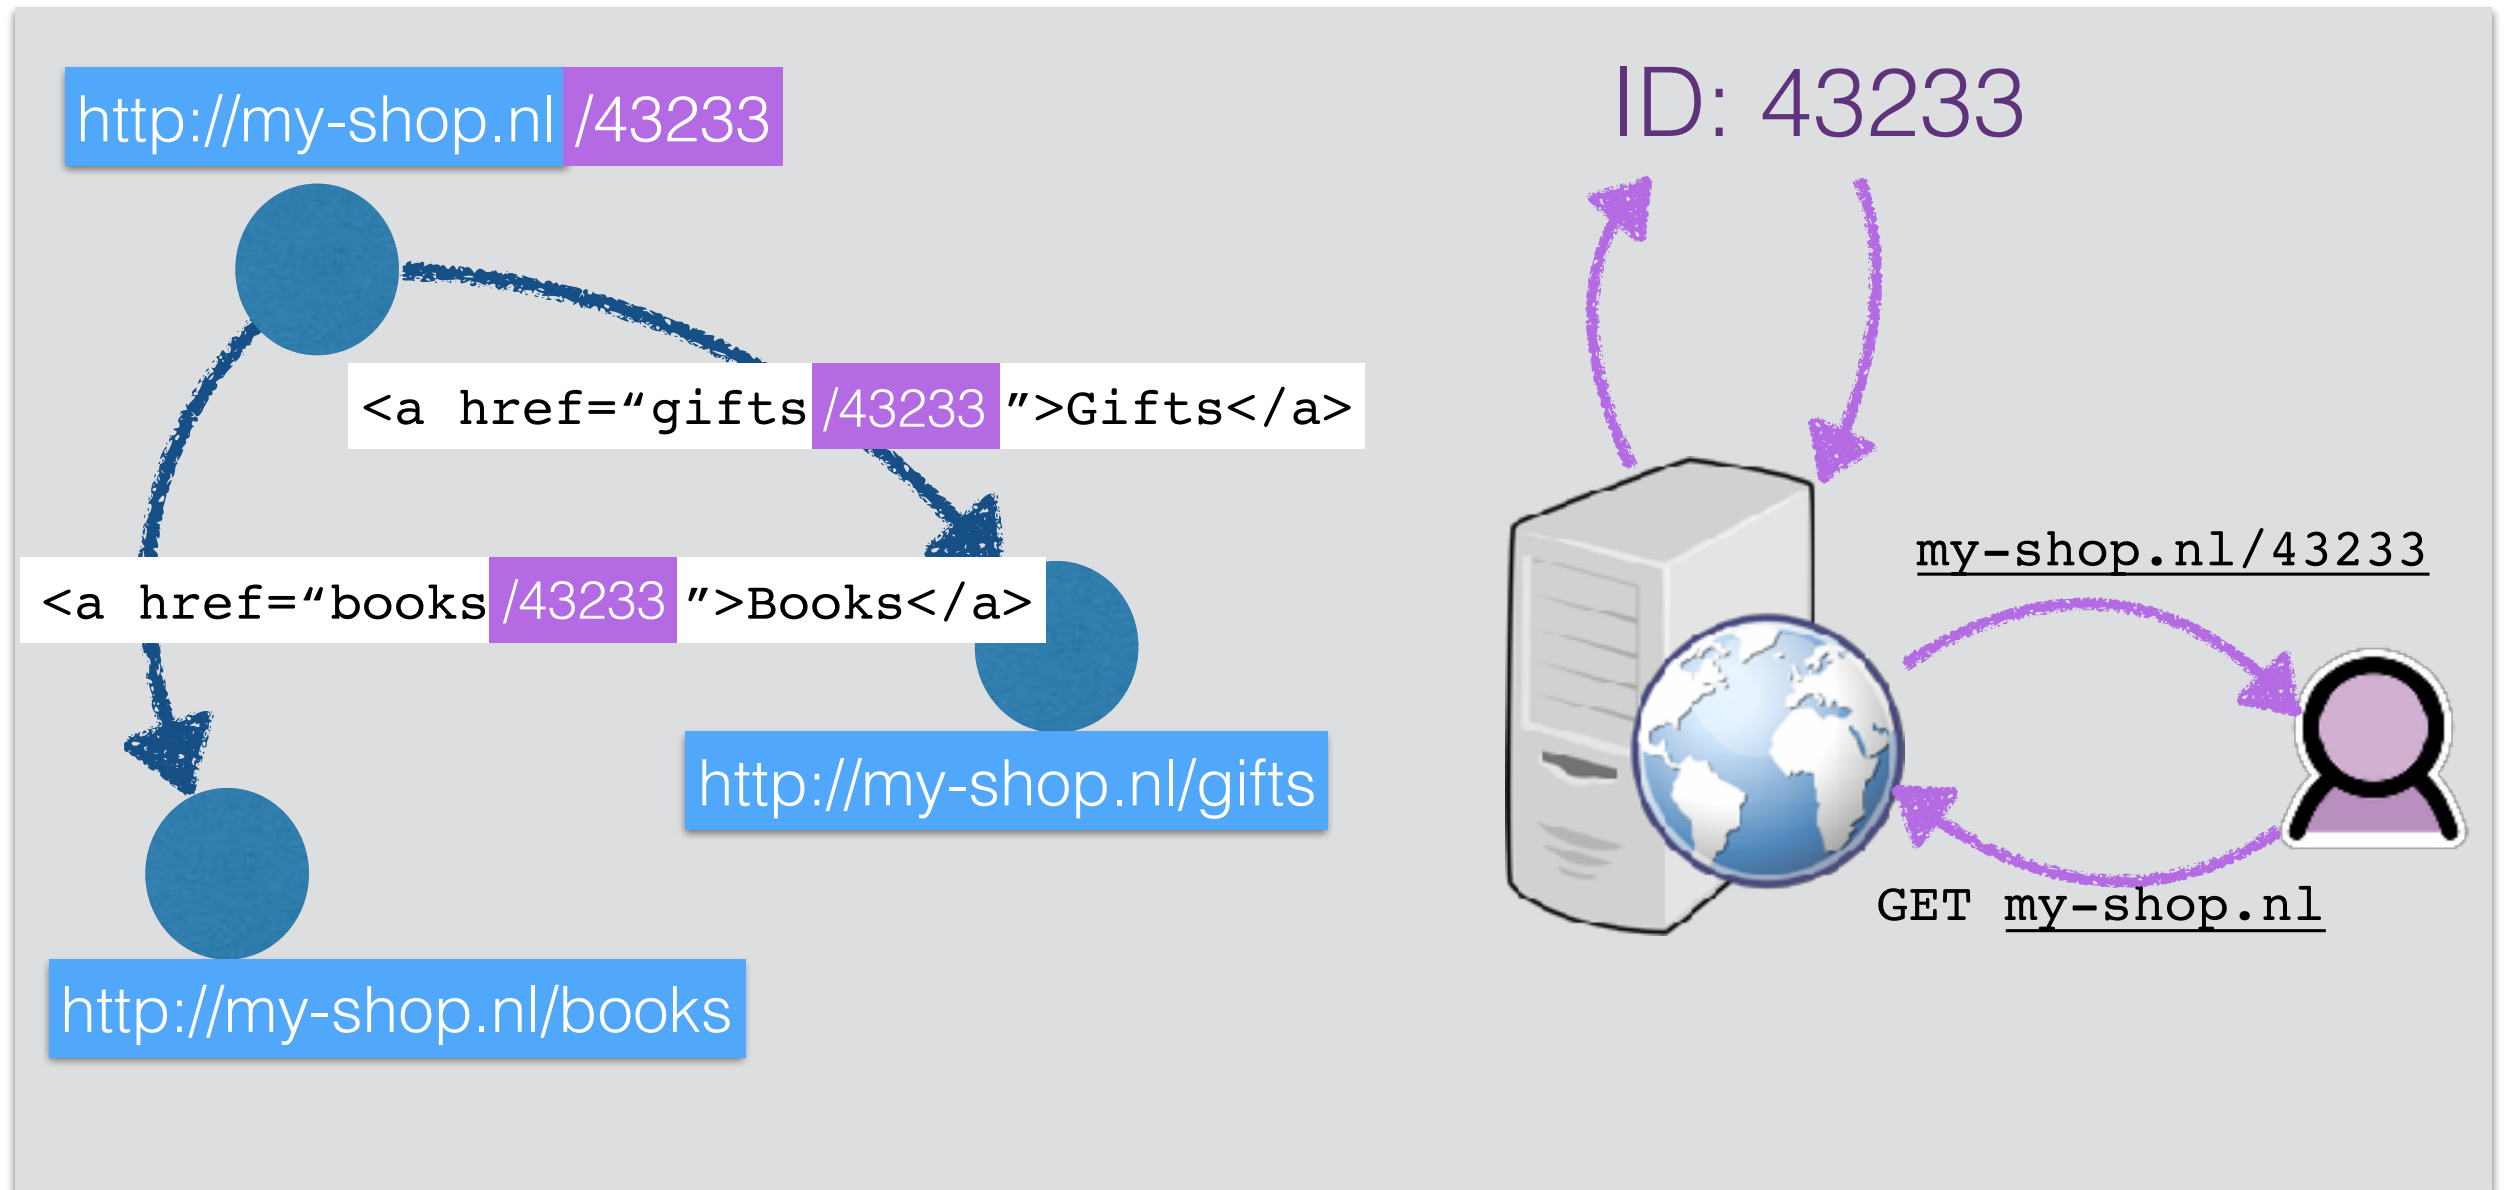
# Fat URLs

- Idea: track users through the **generation of unique URLs**
  - First time a user visits a resource within a Web site, a **unique ID** is generated by the server
  - Server **redirects** client to the fat URL (URL+unique ID)
  - Server **rewrites the HTML** when an HTTP request with a fat URL is received (adds ID to all hyperlinks)

```
<a href="/browse/002-1145265-8016838">Gifts</a>  
<a href="/wishlist/002-1145265-8016838">Wish List</a>
```

- Independent HTTP requests thus tied into a single session

# Fat URLs



# Fat URLs have issues too!

- Fat URLs are **ugly**
- Fat URLs **cannot be shared** (private info shared)
- Fat URLs **break web caching** mechanisms
- **Extra server load** (HTML page rewrites)
- Users can “**escape**” (ID is lost when user navigates away from the site)

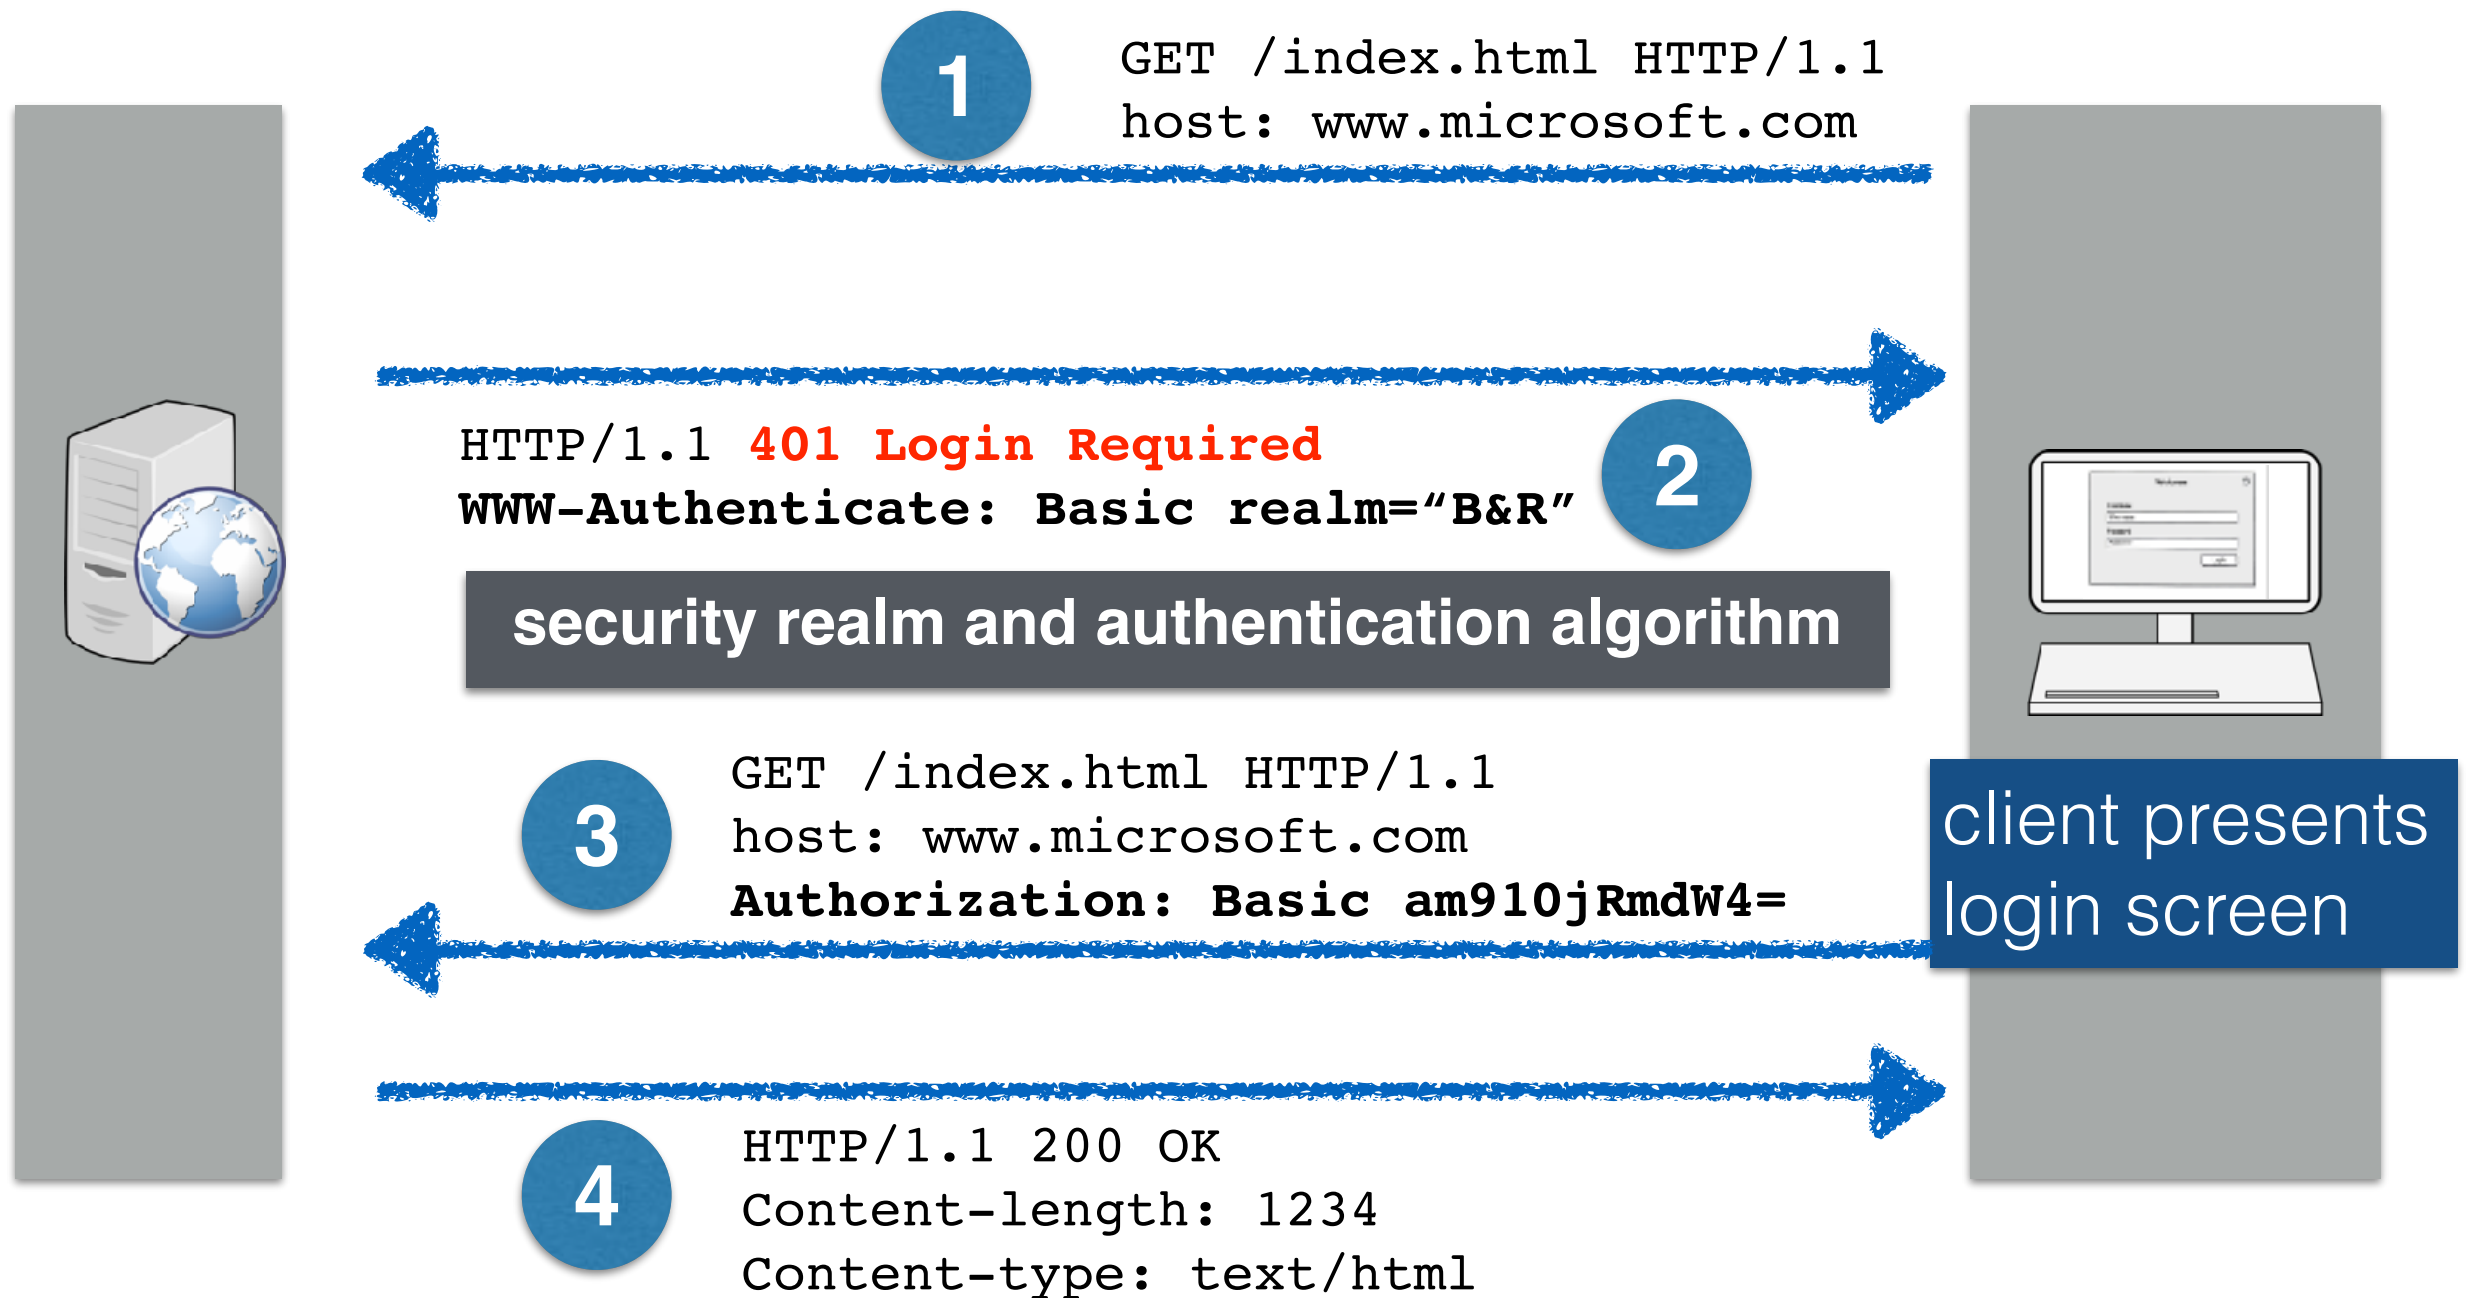


# HTTP basic authentication

- Server explicitly asks the user for authentication (**username and password**)
- HTTP has a **built-in mechanism** to support username/password based authentication via `WWW-Authenticate` and `Authorization` headers
- HTTP is **stateless**: once logged in, the client sends the login information with each request

# HTTP basic authentication

Put into practice  
in Assignment 1



In future HTTP requests to the site, the browser automatically issues the stored username/password

# HTTP basic authentication

Authorization: Basic

- Username and password are joined together by a colon and converted to **base-64 encoding** (binary-to-text encoding)
- Base-64 encoding ensures that **only HTTP compatible characters** are entered into the message (takes as input binary, text and international character data strings)

Normandië	Tm9ybWFuZGnDqw==
Delft	RGVsZnQ=
España	RXNwYcOxYQ==

# HTTP basic authentication: secure?

- Username and password can be **decoded trivially** (sent over the network “in the clear”)
- Users tend to **reuse** login/password combinations; a non-critical web site may use basic authentication without SSL that an opponent can capture and try on critical sites

# HTTP basic authentication: overall

Basic authentication prevents **accidental** or **casual access** by curious users (privacy is desired but not essential).

Basic authentication is useful for **personalisation** and access control within a “friendly” environment (intranet).

“In the wild”, basic authentication should always be used in combination with **secure HTTP (e.g. https)** — avoids sending username/password **in the clear** across the network.

# Secure HTTP

# Secure HTTP

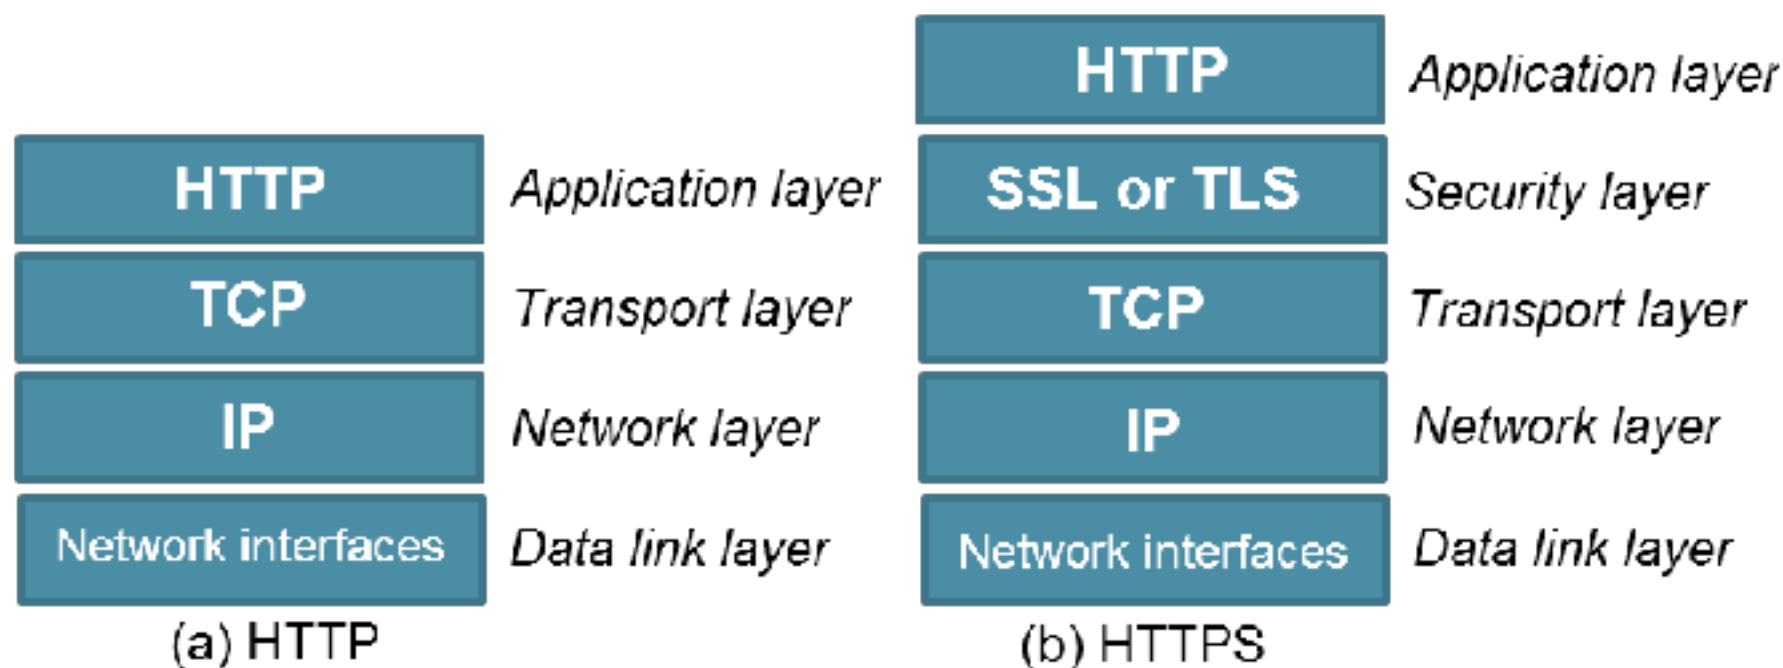
- So far: lightweight authentication
  - Not useful for purchasing, bank transactions or confidential data
- Secure HTTP should provide

- A. Server authentication (client is sure to talk to the right server)
- B. Client authentication (server is sure to talk to the right client)
- C. Integrity (client and server are sure their data is intact)
- D. Encryption
- E. Efficiency
- F. Adaptability (to the current state of the art in encryption)



# Secure HTTP: HTTPS

- HTTPS is the most popular secure form of HTTP
- URL scheme is `https://` instead of `http://`
- Request and response data are **encrypted** before being sent across the network (SSL: Secure Socket Layer)



Client & server **negotiate** the cryptographic protocol to use.